

Lab Report: Localization

Team 21

May Huang
Mohammadou Gningue
Haley Sanchez
Nikhil Kakarla
Neil Chowdhury

Robotics: Science and Systems

April 15, 2023

1 Introduction

1.1 Lab Objective

Our team's objective for this lab was to implement a localization algorithm on our robot. Monte-Carlo Localization (MCL) enables our robot to locate itself within a known environment given sensor data and a motion model, providing information on the robot's position and orientation. A successful, robust, and reliable localization algorithm will allow our robot to locate and navigate through any environment autonomously. Given current pose data and an objective endpoint, the robot can navigate to that point by evaluating the relative error in its current pose and modifying its path of motion. We implemented our MCL algorithm by breaking down into three components: the motion model, the sensor model, and the particle filter.

1.2 Defining Success

In order to define success for this lab, we had a variety of metrics. First, we had unit tests for the motion and sensor model that would allow us to determine if they were working correctly. For the particle filter, we wanted to be able to watch in simulation as the particles followed the robot. Moreover, we expected the particles to follow the robot even when we introduced noise into the system. The main definition of success was that we did not want the particles to drift as the robot continued driving. Finally, we wanted our MCL algorithm to work on the robot. The main indicators of success was that the Lidar scans matched the map of Stata, which would show that the robot has an accurate idea of

its location. Also, we want the standard deviation of our particle locations to be relatively low even in the presence of noise. This indicates that we have an accurate idea of the location of the robot and are not sampling from a wide range of possible locations.

2 Technical Approach

2.1 Developing a Motion Model

2.1.1 Algorithm

In order for our robot to find itself within a given environment, we use a motion model that relies on exteroceptive sensors. This allows for tuning to some amount of uncertainty, in the case that the robot was placed in an unexpected location, orientation, or moved in the middle of driving. Given a distribution of previous particle locations and recent odometry data, our model applies the shifts from the odometry data as well as a rotation matrix to the initial particle positions, then returns an updated pose.

In practice, the robot generates particles, uses the odometry data to assign probability values for each of the particles, evaluates where particles will fall after some displacement or movement of the racecar, and prunes data and the process repeats again.

2.1.2 Noise Function

Random noise is also injected into the system to cause the particles to spread as the robot moves. On every odometry update, we add random noise from a normal distribution from 0 to 0.01 to both position and orientation measurements. This method helps account for inconsistencies in sensor data, for example if the robot drifts or slips while in motion.

2.2 Designing a Sensor Model

2.2.1 Concept

The goal of the sensor model is to be able to assign likelihood weights to hypothesized particles x_k on a map, m , from an observed laser scan, z_k , as seen in Figure 2. Within the overall particle filter process, this allows us to prune the particles based on those probabilities before being re-fed into our motion model in order to create a more accurate pose estimate. The probability of a range measurement is dependent on the linear combination of four different cases for the observed laser scan.

- p_{hit} ; The measurement hit a known obstacle in the map
- p_{short} ; The measurement was too short, hitting an unexpected obstacle or otherwise affected by the environment

- p_{max} ; The measurement was too large, either missing an object or otherwise not returning back to the lidar
- p_{rand} ; The measurement is unexpectedly and randomly incorrect

These cases are linearly combined with four constants α_{hit} , α_{short} , α_{max} , and α_{rand} , to get the correct probability as shown in Figure 1.

$$p(z_k^{(i)} | x_k, m) = \alpha_{hit} \cdot p_{hit}(z_k^{(i)} | x_k, m) + \alpha_{short} \cdot p_{short}(z_k^{(i)} | x_k, m) + \alpha_{max} \cdot p_{max}(z_k^{(i)} | x_k, m) + \alpha_{rand} \cdot p_{rand}(z_k^{(i)} | x_k, m)$$

Figure 1: Equation for Likelihood of a Singular Range Measurement Given an Observation

2.2.2 Algorithm

The sensor model is calculated by first pre-computing a probability table of potential particles and observation measurements from 0 to 200 pixels, representing range measurements in a lidar scan on the map. We pre-compute the table in order to reduce the computation power needed to complete this task and increase the computational efficiency and speed. It also allows us to normalize our probabilities to ensure they sum to 1. Within the table, we normalized probabilities across p_{hit} values and then calculated the likelihood weights according to the equation described in Figure 1. We then normalized across columns in the pre-computed table to make sure that the probabilities across the columns add up to 1. The likelihood of a scan is computed as the product of the likelihoods of each of n range measurements in the scan.

$$p(z_k | x_k, m) = p(z_k^{(1)}, \dots, z_k^{(n)} | x_k, m) = \prod_{i=1}^n p(z_k^{(i)} | x_k, m)$$

Figure 2: Equation for Likelihood of a Scan Given an Observation

After the table is pre-computed, we then evaluate the probabilities of actual sensor measurements and particles. Particles are given in the form of $[x, y, \theta]$ and need to be put through ray casting in order to get the ground truth lidar scan representative of each particle. We then convert values from the lidar scans and ray-casting from meters to pixels by dividing them by `self.map_resolution*lidar_scale_to_map_scale` parameters. These factors were pre-determined based on our environment. We then clipped all lidar and ray-casting scan values greater than 200 to 200, and values less than 0 to 0 in order to be looked up in our pre-computed table. We then used the pre-computed table to calculate the probability of the particles given the measured scan (observation) using the equation in Figure 2. Once, we get the probability from the pre-computed table, we can then squash the

probabilities to reduce variation and peaky-ness by raising them to the power of $\frac{1}{2.2}$.

2.2.3 Down-sampling

We made a design choice in order to further increase our efficiency. Our real-time lidar scan has 1081 measurement points spread of 240° . Many nearby measurement points are redundant and may add unnecessary noise via more frequent peaks in our probabilities. Therefore we downsample to our desired `self.downSampleSize` size of 100 beams. We use interpolation over all 1081 datapoints to get 100 estimated desired range measurements evenly spread across the 240° lidar spread.

2.3 Implementing Particle Filter in Simulation

2.3.1 Algorithm

Our implementation of the particle filter follows the basic structure of most particle filters. Upon an lidar or odometry callback, we call the sensor and motion models respectively. The motion model updates the positions of the particles with the movement of the robot and introduces a baseline level of noise into the system. Additionally, the sensor model compute probabilities for each particle. We then sample from these probabilities and update the particles. This process allows for some error in the particles while still allowing us to update them to match where the robot is. Finally, we compute the new transformation of the robot and publish the particles. The robot pose is calculated using strict mean for the x and y values and the circular mean algorithm for the theta values to handle edge cases. The particle filter model allows us to localize the robot robustly and adjust for error. In fact, the error in the motion model is crucial to the success of the algorithm.

2.3.2 Design Choices

While implementing our particle filter, we made several core design decisions. First, we decided to use a locking mechanism. The motivation for this was that rospi callbacks are not thread safe. This means that certain variables could be changed and accessed by multiple callbacks simultaneously. This was causing our code to crash. Therefore, we used unit locking to ensure that callbacks did not coincide and our program was thread safe. Additionally, we chose to use NumPy operations for all of our code transformations. The reason for this is that NumPy operations are very fast and optimized compared to simple for loops. This decision allowed us to speed up our computation and make a particle filter that can accurately track our robot even when moving at high speeds. Finally, we used caching. Caching is the idea of computing certain values only once and then storing them. We use caching in our sensor model as well as in our particle filter to store data ranges. This again is a core design decision that

speeds up our code and allows us to quickly compute particle updates. This again allows our particle filter to operate at high speeds.

2.3.3 Tuning Noise

When noise was added to the motion model there was no tuning or examples that highly relied on noise. So when the motion model and sensor model came together to implement the particle filter, we needed to re-evaluate the noise scaling factor. We need noise so that the re-sampling process goes smoother and so that the values do not collapse on one another. In figure 2, we experimented with various noise scaling factors. The .001 light grey line shows that with too little noise our absolute error between the ground truth and the noisy odometry is very high. The .1 black line shows that with too much noise the model struggles to keep absolute error low. The .02 scaling value with the blue line shows that this level of noise is able to keep absolute error low as time increases. Based on this data, we used a noise scaling factor of .02 in our particle filter.




Figure 3: Impact of Various Noise Scaling on Absolute Error

2.3.4 Increasing Algorithmic Speed

We navigated a tradeoff between using more particles, which would lend itself to greater locational accuracy, and algorithmic speed. We increased algorithmic speed by downsampling lidar scan data, to where we were working with 100 data points instead of over 1000. We used interpolation to create a smooth curve of the lidar data and then sampled that curve. This process allowed for faster computation on less data points as well as a major reduction in lidar

scan noise. Additionally, our MCL algorithm avoided recomputation by using caches when processing commonly used data, such as scan ranges. Using a precomputed probability table instead of computing probability following every raycast scan greatly increased speed. Without this precomputation, the MCL algorithm would not be able to track the robot when moving at high speeds. Finally, computing using NumPy operations over for loops also helped speed up operations as NumPy operations are very efficient.

2.3.5 Autograder Results

We then tested our algorithm on the autograder. We found that it passed the localization tests for “no odometry noise” and “some odometry noise,” but did not pass the “more odometry noise” test.




Figure 4: No noise




Figure 5: Some noise




Figure 6: More noise

We attributed this sensitivity to noise to our the parameters of our algorithm making it possible for small divergences to drastically effect the rest of the path. We noticed that in the “more noise” case, the robot’s path strayed off at one point and was otherwise correct for the rest of the path, based on the odometry data. We think that better averaging could fix this issue (i.e., by reducing the impact of the angular drift at the point where the paths diverged), or changes to the parameters (i.e., noise added to heading updates). Unfortunately, we did not have sufficient time to attempt changes, as we focused our effort on implementing and tuning localization for the physical robot.

2.4 Localization on Physical Robot

After getting the localization to work in simulation, it was time to transfer the code onto the robot itself and test it in the real world. The transfer process was relatively easy, only requiring us to change a few topic names and update some of the code. The main challenge was that we needed to tweak our locking mechanisms to work on the physical robot, where the hardware runs at different speeds than in simulation. After getting the locking mechanism to work, we were able to eliminate the concurrency issues and the model worked on the robot. We visualized our results in RViz and used rqt plot to gather data about our particle distribution.

2.4.1 Overall Results

Our localization worked very well on the robot. When running it live on the robot, we visualized the data in Rviz. An example output of our data is shown in the below figure. In the output, the base layer is a map of the Stata basement. This is overlaid with the live feed of the lidar scan data. Moreover, we added visualizations for our particles as well as the base link guess of where our robot was. As shown below in figure 7, the model is working because the lidar scans very closely match the walls of the Stata basement in the map. This means that the robot has an accurate idea of where it is on the Stata map and the external readings match the internal location estimate. Driving at different speeds, we were able to show that the lidar scans continually matched the walls around the robot. Therefore, the localization algorithm was functional and effective as the robot moved around its environment.




Figure 7: Functional Output Display in RViz of Localization

2.4.2 Testing

After visually verifying that our algorithm was working, the next step was to run tests to ensure that the localization algorithm was robust and effective under a number of different conditions. First, we tested initializing the robot at different locations. Next, we tested initializing the robot at a different location in the physical environment than where it thought it was in simulation. The results of this experiment are shown below. Our localization algorithm was able to quickly determine that the robot was in the wrong location and adjust the estimation until it was accurate. Below are images of the image outputs at the start of initialization and a few seconds later. In the first image, the lidar scans clearly do not line up with the walls of the Stata map. However, within a few seconds the robot updated its estimated pose and the lidar scans again very closely mirrored the map. This process is shown in the second and third images. This shows that the robot was able to localize and “find” itself in its environment even when the initial guess was incorrect. Additionally, we included a graph in figure 9 of the standard deviation of our particles during this trial. In the graph, the particles first start very scattered with high standard deviation as the robot tries to find itself on the map. Quickly, the robot is able to sample the most correct particles and the standard deviation drops to the baseline level created by our artificial noise. This is proof of our algorithm working, as it is able to sample the correct particles and create a very accurate estimate quickly.




Figure 8: Robot initialized incorrectly and able to localize

Additionally, we wanted to test our localization algorithm with the robot driving at different speeds. Although visually it appeared that the algorithm was working, we conducted another experiment driving the robot first quickly and

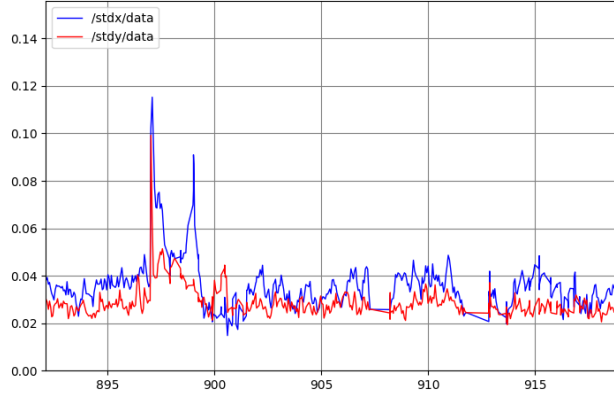


Figure 9: Standard Deviation of x,y coordinates of particles with bad initialization

then slowly. The graph of the standard deviation of the particles is shown below in figure 10. The key result is that, after the initial initialization, the standard deviation did not increase too much and was able to return to its base levels quickly. This is proof that our localization algorithm is robust to driving at different speeds and is computationally quick enough to follow the robot even at its max speed.




Figure 10: Standard Deviation of x,y coordinates of particles while driving

In conclusion, transferring the localization algorithm onto the robot allowed us to verify the accuracy of our algorithm. By testing different initialization conditions and driving speeds, we found that our localization algorithm is correct, robust, and efficient.

3 Conclusion

We achieved our goal of implementing a successful and effective algorithm for MCL localization with our particle filter. We delegated the motion model and sensor model algorithms within our team, which allowed our team to work on them in parallel and to be more efficient. We were able to use unit tests to verify that the motion and sensor models were functioning correctly. Once we put these algorithms to use within the particle filter, we were able to verify its results through the auto grader. In simulation and through unit tests, we then knew that our algorithm was functioning correctly. After that we moved to adapting our algorithm to function on the physical robot. We had to make slight alterations to our original code. Then we completed multiple experiments with the physical robot to verify that the localization was working correctly in Stata basement. From our experiments, we found that our particle filter algorithm functions effectively in the real world.

Next steps include completing further physical experiments on the robot. Another experiment that we could complete is running wall follower and then comparing the actual distance from the wall with the measured distance from the wall with localization. Additionally, we could experiment with various locations if we are able to attain a map of the area like we had for Stata basement.

Overall, we completed implementing our functioning and effective particle filter algorithm that lets us use localization on the robot. So far, we have worked with lidar scans, ZED camera, and implemented localization. We look forward to learning about path finding and implementing it with the physical robot.

4 Lessons Learned

4.1 Haley

Since this lab was started prior to spring break, it was very crucial for us to take notes so that we could pick up right where we left off. Prior to break we had a good conceptual understanding of the lab and how the components mapped to the individual assignments from beforehand. Having good documentation for technical work is super imperative if there is a break or new people joining in on the project. By keeping documentation the project is able to get done smoothly.

4.2 May

This lab was challenging because it required clear and consistent communication between team members while working in parallel with each other. Being able to do the individual pre-assignment was very key for understanding how to approach these algorithms from a conceptual lens, and debugging was a huge part of successfully implementing and testing our code. I think planning ahead for

securing robot time and using it efficiently, as well as communicating successfully with other teams on access to equipment, will be something that continues to be extremely valuable to practice and crucial to finishing work by deadlines.

4.3 Mo

This lab was the most technically challenging so far and taught me the value of both taking good notes and asking questions to make sure the conceptual understanding is clear. When working with others on the team, we learned to fill in each other's technical gaps to ensure we understood and contributed fairly to the work. I enjoyed how technically involved this lab felt.

4.4 Neil

I felt this lab was difficult, but I learned a lot about localization algorithms and implementation details (e.g., NumPy and Cython for efficiency, resampling, circular means). In terms of team working, I learned how division of labor and communication must be reconciled well to accomplish tasks on time. I wrote the motion model and particle filter fairly quickly, but integrating everything with the sensor model took most of the work because I didn't have a complete understanding of how everything would fit together, not having worked on the sensor model. The break also disrupted our work, and I learned that it was important to document code and record thoughts beforehand so it would be easy to pick up where we left off. Overall, though, I felt that this lab was one of the most useful in my understanding of algorithms that are close to real-world applications.

4.5 Nikhil

I learned a lot about this lab. The primary lesson for me was dealing with concurrency issues throughout the code. Because of the structure of ROS, callbacks are not thread safe. Therefore, I had to apply locking to each of the nodes to make sure there were no concurrent transactions that would cause the program to crash. Additionally, I learned a lot about communication and parallel development. For example, our team began to comment code and document our code better. This allowed us to both pick up where we left off more easily as well as code individually. In this way, we were able to streamline our development process and spend less time trying to figure out what other people's code was doing. Therefore, I learned a lot about both technical programming and team communication throughout the course of this lab.

5 Credits Page

5.1 Nikhil

- Localization on Physical Robot: ALL

5.2 Neil

- Motion Model: Noise Function
- Sensor Model: Down-sampling
- Autograder Results

5.3 Mo

- Implementing Particle Filter in Simulation: Algorithm
- Implementing Particle Filter in Simulation: Design Choices
- Autograder Results

5.4 May

- Introduction
- Motion Model: Algorithm
- Implementing Particle Filter in Simulation: Increasing Algorithmic Speed

5.5 Haley

- Sensor Model: Algorithm
- Implementing Particle Filter in Simulation: Tuning Noise
- Conclusion